

How Markov Chain Monte Carlo (MCMC) Works

Markov Chain Monte Carlo (MCMC) is a class of algorithms used to **sample from a probability distribution** for which direct sampling is difficult or impossible. This is often the case for complex, high-dimensional distributions, especially the **posterior distribution** in Bayesian inference.

The core idea behind MCMC is to construct a **Markov chain** – a sequence of random variables where the probability of transitioning to the next state depends only on the current state.¹² This Markov chain is designed such that its **stationary distribution** (the distribution it converges to after a large number of steps) is the target probability distribution we want to sample from.³

Here's a breakdown of the key steps involved:

1. **Define the Target Distribution:** This is the probability distribution we want to sample from (e.g., the posterior distribution in Bayesian inference).⁴ We usually know a function that is proportional to this distribution, even if we don't know the normalization constant. Let's call this unnormalized target distribution $f(x)$. Our goal is to get samples that look like they came from $P(x)=Zf(x)$, where $Z=\int f(x)dx$ (the normalization constant) might be unknown.
2. **Initialize:** Start at an arbitrary initial state x_0 in the sample space of the target distribution.⁵
3. **Propose a New State:** From the current state x_t , propose a new candidate state x^* according to a **proposal distribution** $Q(x^*|x_t)$.⁶ The choice of the proposal distribution is crucial for the efficiency of the MCMC algorithm.⁷ Common choices include:
 - **Random Walk:** Proposing a new state by taking a random step from the current state (e.g., adding a random value from a normal distribution).
 - **Independent Proposal:** Proposing a new state independently of the current state (though this requires some knowledge of the target distribution to be efficient).
4. **Accept or Reject the Proposed State:** Decide whether to move to the proposed state x^* or stay at the current state x_t .⁸ This decision is made based on an **acceptance probability** $\alpha(x^*|x_t)$.⁹ The acceptance probability is designed to ensure that the Markov chain converges to the target distribution. A common acceptance criterion is the **Metropolis-Hastings algorithm**:
$$\alpha(x^*|x_t) = \min(1, \frac{f(x^*)Q(x_t|x^*)}{f(x_t)Q(x^*|x_t)})$$
 - The ratio $\frac{f(x^*)}{f(x_t)}$ compares the "likelihood" of the proposed state to the current state under the target distribution.
 - The ratio $\frac{Q(x_t|x^*)}{Q(x^*|x_t)}$ accounts for the asymmetry in the proposal distribution (how likely it is to propose x^* from x_t versus proposing x_t from x^*). If the proposal distribution is symmetric (i.e., $Q(x^*|x_t) = Q(x_t|x^*)$), the acceptance probability simplifies to:
$$\alpha(x^*|x_t) = \min(1, \frac{f(x^*)}{f(x_t)})$$
5. **Update the State:**
 - Generate a random number u from a uniform distribution between 0 and 1 ($U(0,1)$).¹⁰
 - If $u \leq \alpha(x^*|x_t)$, then **accept** the proposed state and set $x_{t+1} = x^*$.
 - If $u > \alpha(x^*|x_t)$, then **reject** the proposed state and stay at the current state, setting $x_{t+1} = x_t$.
6. **Repeat:** Continue steps 3-5 for a large number of iterations.¹¹
7. **Collect Samples:** After a "burn-in" period (initial iterations discarded to allow the chain to converge to the stationary distribution), the sequence of accepted states ($x_{\text{burn-in}+1}, x_{\text{burn-in}+2}, \dots, x_T$) forms a sample from the target distribution.

Why does this work?

The acceptance probability is carefully constructed to satisfy the **detailed balance condition** for the Markov chain with respect to the target distribution. Detailed balance ensures that the rate of transitioning from state x to state y is

equal to the rate of transitioning from state y to state x under the target distribution. If detailed balance holds, then the target distribution is the stationary distribution of the Markov chain.¹²

Example: Sampling from a Beta Distribution using Metropolis-Hastings

Let's say we want to sample from a Beta distribution with parameters $\alpha=2$ and $\beta=5$. The probability density function (PDF) of a Beta distribution is:

$$P(x|\alpha,\beta)=B(\alpha,\beta)x^{\alpha-1}(1-x)^{\beta-1}, \text{ for } 0 \leq x \leq 1$$

where $B(\alpha,\beta)$ is the Beta function (the normalization constant). Let $f(x)=x^{\alpha-1}(1-x)^{\beta-1}$ be the unnormalized target distribution.

We will use a simple **random walk proposal** where we propose a new state x^* by taking a step from the current state x_t :

$$x^*=x_t+\epsilon$$

where ϵ is drawn from a symmetric distribution centered at zero, such as a uniform distribution $[-0.1,0.1]$. Since the proposal distribution is symmetric ($Q(x^*|x_t)=Q(x_t|x^*)$), the acceptance probability simplifies to:

$$\alpha(x^*|x_t)=\min(1, f(x_t)/f(x^*))=\min(1, (x_t)^{2-1}(1-x_t)^{5-1}/(x^*)^{2-1}(1-x^*)^{5-1})=\min(1, x_t(1-x_t)/x^*(1-x^*))$$

Steps:

1. **Initialize:** Let's start with $x_0=0.5$.
2. **Propose:** Draw ϵ from $U(-0.1,0.1)$. Let's say $\epsilon=0.05$. Then $x^*=0.5+0.05=0.55$.
3. Calculate Acceptance Probability:
 $f(x_t)=0.5(1-0.5)^4=0.5 \times 0.0625=0.03125$
 $f(x^*)=0.55(1-0.55)^4=0.55 \times (0.45)^4=0.55 \times 0.04100625 \approx 0.02255$
 $\alpha(0.55|0.5)=\min(1, 0.03125/0.02255)=\min(1, 0.7216)=0.7216$
4. **Accept or Reject:** Generate a random number u from $U(0,1)$. Let's say $u=0.3$. Since $u=0.3 \leq 0.7216$, we **accept** the proposed state. So, $x_1=0.55$.
5. **Repeat:** We continue this process for many iterations. The sequence of accepted x values will eventually approximate samples from the Beta(2, 5) distribution.

Important Considerations:

- **Burn-in Period:** The initial samples of the Markov chain might be far from the stationary distribution, so they are typically discarded.¹³
- **Mixing:** The rate at which the Markov chain converges to the stationary distribution is crucial. A chain that mixes slowly will require many more iterations to produce representative samples.
- **Proposal Distribution:** The choice of the proposal distribution significantly affects the efficiency of the MCMC algorithm.¹⁴ A good proposal distribution will lead to a reasonable acceptance rate and good mixing.¹⁵
- **Autocorrelation:** Successive samples in an MCMC chain are often correlated.¹⁶ Techniques like thinning (keeping only every n -th sample) can be used to reduce autocorrelation.¹⁷

MCMC algorithms are powerful tools for sampling from complex distributions and are widely used in Bayesian statistics, computational physics, and other fields where direct sampling is intractable.¹⁸ Common MCMC algorithms include Metropolis-Hastings, Gibbs sampling, and Hamiltonian Monte Carlo.